



Collaborative filtering

[Learning steps](#)

[Resources](#)




[Quick notes](#)

[YouTube video \(→ \[link\]\(#\)\)](#)

Learning steps

- ✓ [yt-video](#)
 - ✓ [notebook/own-implementation](#)
 - ✓ [book-chapter](#)
-

Resources

- website 
 - [lesson 7](#)
 - notebooks 
 - [Collaborative Filtering Deep Dive](#)
 - Road to the top: [part 3](#) and [part 4](#)
 - book 
 - [chapter 8](#)
 - [solutions to exercises](#)
-

Quick notes

YouTube video (→ [link](#))

- we usually tweak the first or last layers
 - we'll take a look at the rest later on
- road to the top part 2 (RTTT 2)-
 - ConvNext model used to start with + some pre-processing (TTA)
- RTTT 3
 - we'll use larger models ⇒ more parameters
 - they can find more patterns etc
 - the gradients are more numerable and take up more computation space on the GPU

- potential memory issues when training a model on a GPU (*CUDA out of memory* e.g.)
 - how much memory will a model use? → use `func report_gpu()` to find out
 - to solve it: use `GradientAccumulation`
 - we define a batch size to be divided by a *accum* number
 - we then consider a training loop (compute loss, backward on it, subtract, reset)
 - sometimes, we change the training loop: the gradients are accumulated instead of being reset each time



The difference between GPU classes (RTX3070 Ti, 3080, etc) is not the performance but the memory size. On a smaller GPU mem, use smaller batch size and gradient accumulation!

- rule of thumb for batch size and learning rate
 - batch size: as big as possible. can use multiples of 8 or powers of 2
 - learning rate: if batch size divided by 2, same for LR (not always perfect)
- transformers use a fixed image size, make sure to resize the images to squares of this size
- RTTT 4
 - we'll now work on the last layer of the neural net
 - we'll use a `DataBlock` (one level deeper ⇒ more flexibility)
 - in Pandas, you can specify an index column for a dataframe → you can then use this df as a dictionary
 - we want a model that predicts 20 things: the 10 diseases and the 2 varieties
 - cross-entropy loss
 - softmax
 - the model must say which one of the available categories it is ⇒ must choose one, no place for uncertainty because it adds up to 1
 - we obtain the prob for each category as an output
 - target value (one-hot encoded)
 - then, we sum up across all categories and multiply the target value by the log of the predicted prob (result of softmax)
 - we compute it for every row and add it up
- collaborative filtering deep dive (pretty much chap 8 of the book)
 - goal: recommendation system for movies based on user preferences (ratings)
 - we don't have info about people preferences, so we'll use SGD to find it out (in excel)
 - we decide (randomly) that we have important five factors (parameters instantiated randomly) for each user and for each movie

- we make the product by matrix multiplication to create random predictions for each user for each movie \Rightarrow we need to optimize the parameters
- we use root mean square error (RMSE)